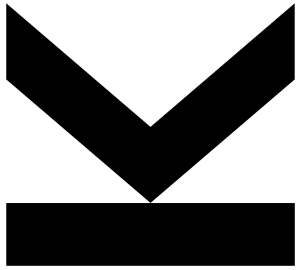


# Assembly Functions



Systems Programming

# What is a „function“?

## ■ Name

- ☐ Symbol representing the address where the function begins

## ■ Parameters

- ☐ Input; arbitrary number and type should be supported
- ☐ Sometimes also output or input+output parameters desirable

## ■ Local variables

- Often called “automatic variables”
- ☐ Private data storage
- ☐ Not accessible from outside
- ☐ “Thrown away” when function ends

## ■ Global variables

- ☐ Public data storage
- ☐ Accessible from inside and outside

# What is a „function“?

## ■ Return address

- ☐ Invisible “parameter”
- ☐ Tells program where to resume executing after function completed
- ☐ In most progr. languages, return address is handled automatically
  - Even in most assembly languages, e.g. X86-\*

## ■ Return value

- ☐ Transfer result back to caller
- ☐ Most programming languages allow only a **single** return value
  - Most also require that it fits in a **single register** (or use one of the circumventions listed below themselves to emulate this)
- ☐ More than one/larger return value needed?
  1. Use pointers to data (i.e. memory addresses of data)
  2. In the function change/set (parameter) values that a pointers point to
  3. Read values in caller after function returns
  - Pointer to data as direct return value
    - Careful: Who reserved the memory, who frees it?
- ☐ Can be “extended” through in+out parameters

# Stack

- Where do we store return addresses, parameters, local variables...?
  - Sometime also large return values (structures)
- In a special memory area called the “Stack”
- We don’t know how “deep” functions will be called (recursion!), so we cannot set an upper limit for the stack size
  - Practically, most OS **do** set a limit → if reached, the program is terminated (to prevent e.g. runaway recursion). But until then only as much memory is provided as is actually used!
- Dangerous from the security point of view
  - Attackers can manipulate it: arbitrary data at arbitrary position (within the stack)
  - Might be executed (but see modern precautions!)

# Stack

- Region at the “top” addresses of memory (of current process)
  - Stack is separate for each process
  - Multiple threads: separate stack for each thread
    - Obviously not at the very top any more, as threads share memory, so stack size limits are more stringent (relocating a stack is impossible!)
    - On 64 Bit computers the logical address space is very large, so this is less a problem
- Stack **grows down** in memory (from high towards low addresses)
- Limited in size: Everything larger than a few kB should be put on Heap
- Used to **implement functions**
  - Save state of the caller
  - Pass parameters to the callee
  - Store return address
  - Store local variables
  - Store large return data (reserved by caller)

# Stack

- **RSP register** always points to the “top” of the stack
  - = **Lowest** used address
    - For a quadword (=8 bytes) it is the “first” byte (=lowest address of the 8)
- **Push data (push)**
  - “Decrement” RSP register
    - How much? As many bytes as the new data item is long!
  - Write data item at new top
- **Pop data (pop)**
  - Read data from top (and store it in some place, typ. some register)
  - “Increment” RSP register
    - How much? As many bytes as we remove – This need **NOT** be the same as was used on push!
- RSP can also be modified by **SUB/ADD**, e.g. for creating/destroying local variables (remember: stack grows down, so **SUB creates** space!)
  - Then the “new content” is **NOT** initialized (=old values)!

# Calling conventions

## ■ How to call a functions:


- ☐ How do we pass parameters? Any metadata (=type) for them?
  - Which registers? Or stack only?
  - Object-oriented languages: Where do we put the “this” pointer?
- ☐ What about local variables?
  - And who cleans up after them?
- ☐ What about the return value (“A” register, stack...)?
- ☐ Which register may be used in the function (=who saves them)?
- ☐ Return address (x86 → hardware restrictions!), frames...

## ■ This is a „calling convention“

- ☐ „Convention“ because there are few technical restrictions
- ☐ Every programming language (or programmer) can decide on his own what to do/how to do it
- ☐ You can even mix them in a single program
- ☐ BUT: However a function is programmed, this function must be called in an exactly matching way!

# Calling conventions: cdecl

- C declaration → Original C programming language
  - All parameters are passed on the stack
  - Stored in reverse order (last parameter is pushed first)
  - Return value in A (=EAX) register
  - Registers A, C, D (=EAX, ECX, EDX) are caller saved, rest is callee saved
  - EBP register used for frame pointer
  - Caller cleans up stack
  - Linux modification: Stack must be 16-Byte aligned on function call
  - Used in Linux-x86-32 Bit
  - Typical C declaration: `void _cdecl funct () ;`



Specify the calling convention to use for this function



# Calling conventions: Others

- Pascal: Pascal programming language
  - ☐ Similar to cdecl, but parameters are pushed on the stack in normal order, which prevents functions with variable count of parameters
  - ☐ Callee has to clean up stack
  - ☐ Used by Windows 16 Bit (Windows 3.x)
- Stdcall: Similar to Pascal
  - ☐ Parameters are pushed in reverse order, like in cdecl
  - ☐ Standard calling convention for Windows 32 Bit
- Microsoft fastcall: `__fastcall`
  - ☐ First two argument are passed in ECX and EDX, rest on stack in reverse order
  - ☐ Windows 32 Bit (depending on compiler; used for optimization)

# Calling conventions: Microsoft X64

- RCX, RDX, R8, R9 are used for the first four parameters, the rest are pushed on the stack in reverse order
  - Smaller values are right-justified in registers (=lower bits)
- Return value is in RAX (smaller values do **not** set upper bits to 0!)
- RAX, RCX, RDX, R8, R9, R10, R11 are caller-saved
- RBX, RBP, RDI, RSI, RSP, R12, R13, R14, R15 are callee-saved
- Caller must always reserve 32 Byte of space on stack (shadow space) for the first four parameters (even if less used!)
  - Note that only space is provided, the parameter values are **not** written there by the caller – they are **only** in the registers!
- Stack pointer must be aligned to 16 Bytes
- Used on Windows 64 Bit

# Calling conv.: **SystemV AMD64 ABI**

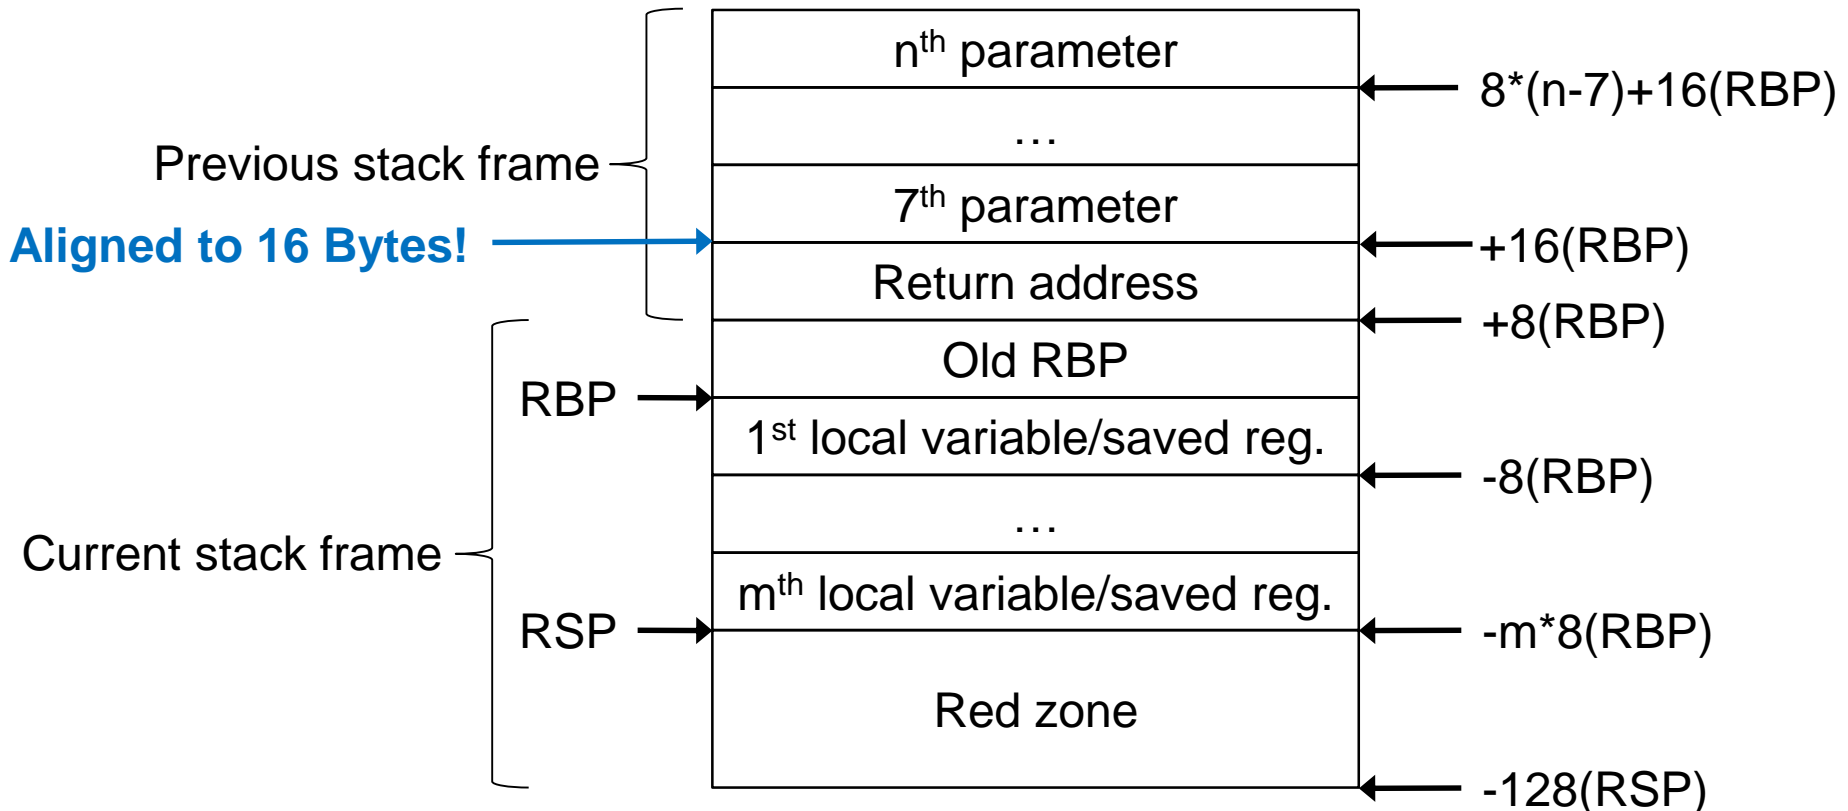
- SystemV (old Unix version; 1983), AMD64 (“original” 64 Bit version of IA32), ABI (“Application Binary Interface” ≈ calling convention+...)
- First six parameters are passed in RDI, RSI, RDX, RCX, R8, and R9, the rest are pushed on the stack in reverse order
  - Fewer parameters? Function can use them freely (caller-saved!)
  - Linux Kernel calls: RCX replaced by R10; max. 6 parameters; no stack ever used (except as memory for where pointers point to)
- Return value is store in RAX (+ potentially RDX for long values)
- RAX, RCX, RDX, RSI, RDI, R1-R11 are caller-saved
- RBP, RBX, R12-R15 are callee-saved

# Calling conv.: SystemV AMD64 ABI

- Stack pointer must be aligned to 16 Bytes (=end of parameters)
- 128 Bytes below RSP are guaranteed to exist and can be used by function for e.g. local variables without RSP adjustment (“Red zone”)
  - Not used by signals and interrupt handlers
  - Will be overwritten by function calls → useful for “leaf” functions
    - “Leaf” function = function that doesn’t call **any** other function
  - Optimization purpose: Use without adjusting RSP
- Used on **Linux 64 Bit**, MacOS...

# Stack frame

- How does the stack look like when a function is called?
  - For each function (without optimizations used!) that is called, a “frame” is reserved



# Calling a function – Caller

- Caller adjusts stack pointer so it will end up at 16-Byte alignment after the next 3 bullets
- Caller saves (typically on stack) all caller-save registers - **if needed**
- Caller pushes parameters N to 7 on stack in reverse order
- Caller puts parameters 1 to 6 in the appropriate registers
- Caller executes the `call` . . . instruction
  - Which pushes address of the next instruction (=RIP) on the stack
    - This is the “return address”, where execution resumes after the function returns
  - The RIP register is modified to contain the address specified in the call instruction – which is the start of the function

↑ Can be exchanged ↓

# Calling a function – Callee - Prologue

- Callee saves the base pointer on the stack
  - “Trace back” to the previous function; used e.g. by debuggers
  - `push %rbp`
- Callee copies the stack pointer into the base pointer
  - `movq %rsp, %rbp`
- Callee subtracts amount of bytes needed for local variables from RSP
  - `subq $???, %rsp`
    - If less than 128 needed (with registers!), the red zone can be used
- Callee pushes all callee-save registers to be used on the stack
  
- Now the actual function begins
  - Parameters: RBP+16 (7<sup>th</sup> parameter) and up
  - Local variables: RBP-8 (1<sup>st</sup> local variable) and down
  - Callee stores return value in RAX (if function and not procedure)

# Calling a function – Callee - Epilogue

- Callee restores all callee-saved registers
- Callee resets the stack to remove the local variables
  - ☐ Should be done even if the red zone was used
  - ☐ Sometimes used to restore stack in case of exceptions
    - Scanning the code for the epilogue to correctly unwind it
  - ☐ `movq %rbp, %rsp`
- Callee restores the old frame pointer
  - ☐ `popq %rbp`
- Caller returns to calling program
  - ☐ `ret`



# Calling a function – Back at Caller

- All local variables have been destroyed
  - **Never** attempt to return a pointer to a local variable!
- Future stack pushes will overwrite the values
  - Some might still be accessible, as the red zone of the current stack frame might cover them (partially)
  - Red zone: Purely an optimization
    - Need not set RBP to RSP and need not subtract from RSP
    - **RSP** is used as base for all parameters and local variables
    - RBP is no longer needed and can be used as normal register
    - Can be used for temporary data and local variables
    - **Careful! Further function calls start from RSP and NOT from the end of the red zone, so calling a function will destroy data in it!**
- Return value is in RAX (plus potentially RDX)
- Caller has to remove parameters 7 to N from stack
  - `addq $(N-6)*8, %rsp` or N-6 times `popq %<some register>`
- Caller restores any saved caller-save registers – **if done**
- Caller re-adjusts stack if any adjustment for alignment was made (or ignores this and just wastes the space until it is cleaned up when it returns itself and resets the stack)

↕ Can be exchanged

# THANK YOU FOR YOUR ATTENTION!

**Slides by: Michael Sonntag**

michael.sonntag@ins.jku.at

+43 (732) 2468 - 4137

S3 235 (Science park 3, 2<sup>nd</sup> floor)



JOHANNES KEPLER  
UNIVERSITÄT LINZ



INSTITUTE  
OF NETWORKS  
AND SECURITY

<https://www.ins.jku.at>



**JOHANNES KEPLER  
UNIVERSITÄT LINZ**

Altenberger Straße 69  
4040 Linz, Österreich  
[www.jku.at](http://www.jku.at)