

# Assembly: Function calls



Systems Programming

# Calling convention: SystemV AMD64 ABI

## ■ About the name

- ☐ SystemV: old Unix version (1983), AMD64: “original” 64 Bit version of IA32
- ☐ ABI (“Application Binary Interface”): The standards that programs must follow to be interoperable with each other and with the system that they run on (e.g. calling convention, etc.)

## ■ Used on 64-bit Linux, etc.

## ■ Function parameters/arguments

- ☐ First six are passed in registers: RDI, RSI, RDX, RCX, R8, R9
- ☐ Further parameters are pushed onto the stack in reverse order
- ☐ Linux Kernel calls: RCX replaced by R10; max. 6 parameters; no stack ever used (except as memory for where pointers point to)

## ■ Return value

- ☐ Stored in RAX (+ potentially RDX for long values)

## ■ Register use

- ☐ RAX, RCX, RDX, RSI, RDI, R1-R11 are caller-saved
- ☐ RBP, RBX, R12-R15 are callee-saved

# Program start

- When Linux starts a program, how does the CPU content look like, where does it begin, how are parameters provided...?
  - Do not assume specific content of registers, unless noted below
    - Flags do have defined content, but for security set them explicitly
  - RSP points to the end of the stack
    - (%RSP) → Number of arguments
    - 8(%RSP) → Pointer to first argument (= program name)
      - Note: This is a **pointer**. This is **not** the string, but the **address** of the **first character** of the string!
    - 16(%RSP) → Pointer to second argument (= first parameter), if present
    - Higher on stack: More parameters, process environment and other data
      - Not used in this course!
  - RDX: Function pointer to specify an exit procedure
    - Not used in this course! Simply ignore it (and use the register)
  - Program entry point: “\_start”
    - Exactly this name, cannot be changed

# Calling a function – Example

## ■ We will call the following function:

- `int doSomething(int p1, int p2, int p3,  
                  int p4, int p5, int p6,  
                  int p7, int p8, int p9)`

- 9 parameters: p1, p2, ..., p9 (64-bit integer each)

- 1 return value (64-bit integer)

## ■ Example for calling this function:

- `if (doSomething(1, 2, 3, 4, 5, 6, 7, 8, 9) != 0) { ... }`

- Calling the function puts the return address on the stack

- Caller is responsible for passing parameters in the right place

# Calling a function – Example

## ■ We will call the following function:

```
int doSomething(int p1, int p2, int p3,  
               int p4, int p5, int p6,  
               int p7, int p8, int p9)
```

- 9 parameters: p1, p2, ..., p9 (64-bit integer each)
- 1 return value (64-bit integer)

## ■ Example for calling this function:

- `if (doSomething(1, 2, 3, 4, 5, 6, 7, 8, 9) != 0) { ... }`
- Calling the function puts the return address on the stack
- Caller is responsible for passing parameters in the right place

The first 6 parameters are stored in **registers**:

- |            |            |
|------------|------------|
| ■ p1 → RDI | ■ p4 → RCX |
| ■ p2 → RSI | ■ p5 → R8  |
| ■ p3 → RDX | ■ p6 → R9  |

# Calling a function – Example

## ■ We will call the following function:

```
□ int doSomething(int p1, int p2, int p3,  
                  int p4, int p5, int p6,  
                  int p7, int p8, int p9)
```

- 9 parameters: p1, p2, ..., p9 (64-bit integer each)
- 1 return value (64-bit integer)

## ■ Example for calling this function:


- `if (doSomething(1, 2, 3, 4, 5, 6, 7, 8, 9) != 0) { ... }`
- Calling the function puts the return address on the stack
- Caller is responsible for passing parameters in the right place

Further parameters are pushed onto the **stack** in reverse order:

(before storing the parameters) RSP

(after storing the parameters) RSP

Stack:



`pushq #p9#`

`pushq #p8#`

`pushq #p7#`

# Calling a function – Example

## ■ We will call the following function:

- `int` doSomething(int p1, int p2, int p3,  
int p4, int p5, int p6,  
int p7, int p8, int p9)

- 9 parameters: p1, p2, ..., p9 (64-bit integer each)

- 1 return value (64-bit integer)

## ■ Example for calling this function:

- `if (doSomething(1, 2, 3, 4, 5, 6, 7, 8, 9) != 0) { ... }`

- Calling the function puts the return address on the stack

- Caller is responsible for passing parameters in the right place

Return value will be stored in  
**register RAX** (by the function)

# Calling a function – Example

## ■ We will call the following function:

```
□ int doSomething(int p1, int p2, int p3,  
                  int p4, int p5, int p6,  
                  int p7, int p8, int p9)
```

□ 9 parameters: p1, p2, ..., p9 (64-bit integer each)


□ 1 return value (64-bit integer)

## ■ Example for calling this function:

```
□ if (doSomething(1, 2, 3, 4, 5, 6, 7, 8, 9) != 0) { ... }
```

□ Calling the function puts the return address on the **stack**

□ Caller is responsible for passing parameters in the right place





# Calling a function – Example

- Internally, the function will also need:
  - Registers: RBX, R10, R11, R12
  - Two 8-byte values as local variables

# Calling a function – Example

■ Internally, the function will also need:

- Registers: **RBX**, R10, R11, **R12**
- Two 8-byte values as local variables



**Callee-saved**

- Caller does not need to do anything
- Must be preserved by the function itself (see later)

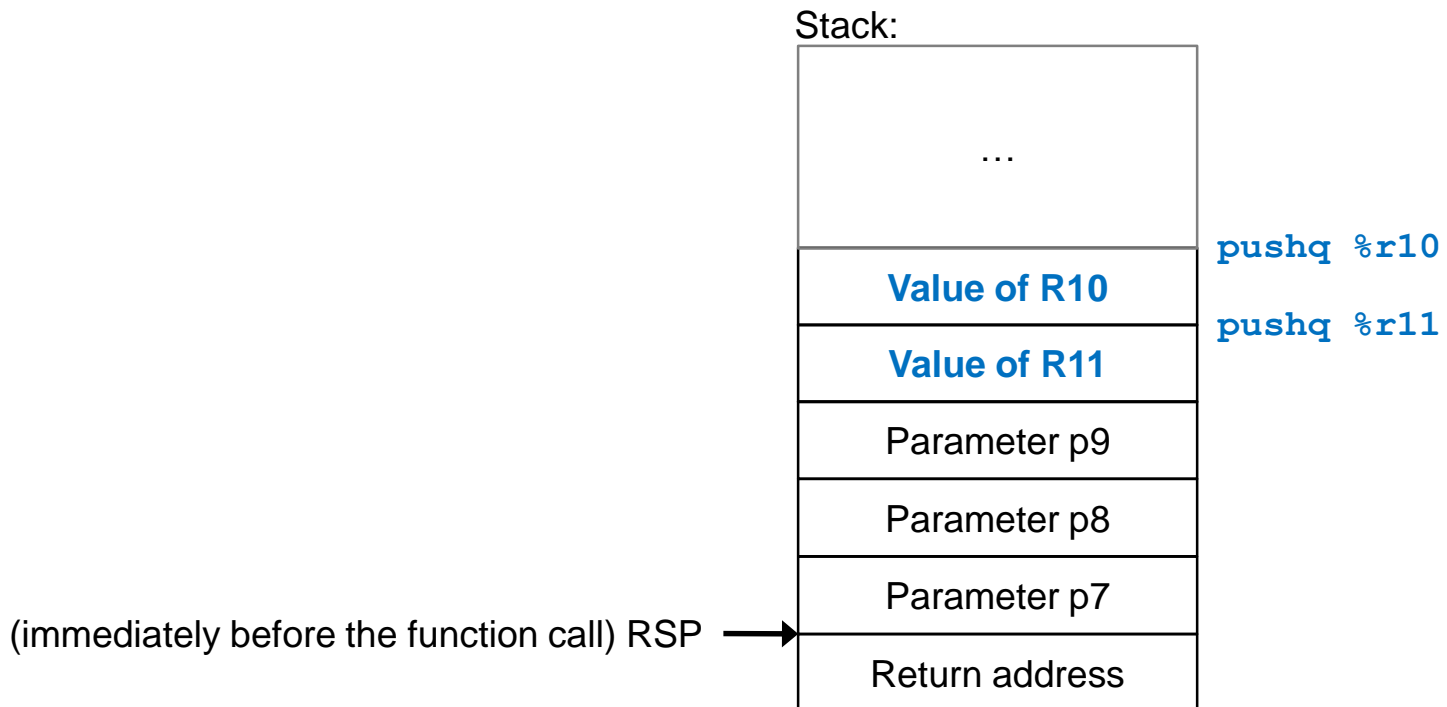
# Calling a function – Example

■ Internally, the function will also need:

- Registers: RBX, R10, R11, R12
- Two 8-byte values as local variables

**Caller-saved**

→ Caller must store these values on **stack**:



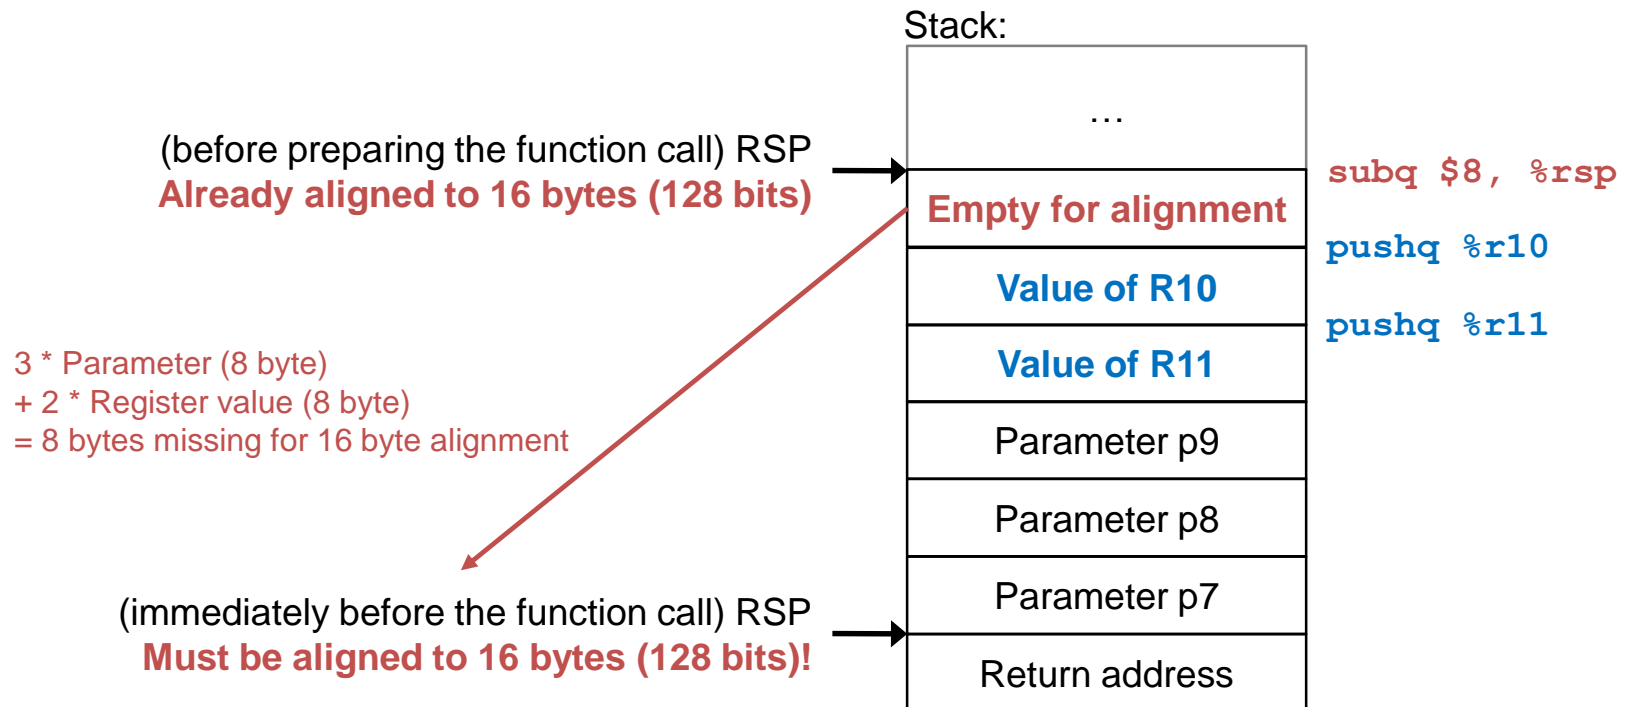
# Calling a function – Example

■ Internally, the function will also need:

- Registers: RBX, R10, R11, R12
- Two 8-byte values as local variables

**Caller-saved**

→ Caller must store these values on **stack**:



# Calling a function – Example

■ Internally, the function will also need:

- ☐ Registers: RBX, R10, R11, R12
- ☐ Two 8-byte values as local variables



→ The function is responsible for this (see later)

# Calling a function – Example (caller)

alignment	{	<code>subq \$8,%rsp</code>	# Ensure stack alignment (we push 24 bytes; if # aligned before we need to "add" 8 bytes more)
caller-saved registers	{	<code>pushq %r10</code> <code>pushq %r11</code>	# Save caller-safe registers
parameters in registers	{	<code>movq \$1,%rdi</code> <code>movq \$2,%rsi</code> <code>movq \$3,%rdx</code> <code>movq \$4,%rcx</code> <code>movq \$5,%r8</code> <code>movq \$6,%r9</code>	# Store first parameter in register # Note: No parameters names appear in assembler! # Store sixth parameter in register
parameters on stack	{	<code>pushq \$9</code> <code>pushq \$8</code> <code>pushq \$7</code>	# Further parameters are pushed on stack # in reverse order!
		<code>call doSomething</code>	
cleanup	{	<code>addq \$24,%rsp</code> <code>popq %r11</code> <code>popq %r10</code> <code>addq \$8,%rsp</code> <code>cmpq \$0,%rax</code> <code>je ...</code>	# Clean up parameters from stack (equal to 3*popq) # Restore caller-safe registers # Clean up alignment space # Now check the return value # If zero, jump over the next block

- Note: We have to do the alignment at the beginning
- Or we would not know where exactly parameter 7 is on the stack!

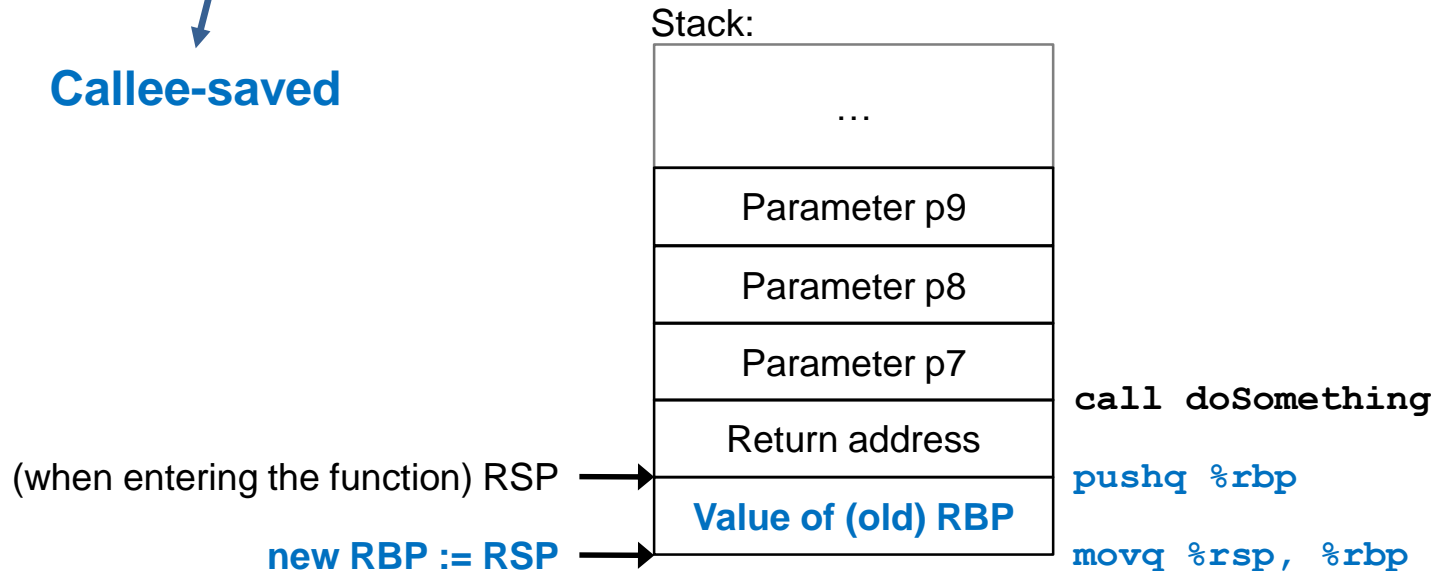
# Calling a function – Example

- The function needs to access its parameters and variables on the stack
  - Stack pointer changes when pushing to/popping from the stack
    - ➔ Cannot be used (or only with lots of difficulties ➔ Compilers do this)
  - Base pointer RBP is used to store that stack position

# Calling a function – Example

- The function needs to access its parameters and variables on the stack
  - Stack pointer changes when pushing to/popping from the stack
    - Cannot be used (or only with lots of difficulties → Compilers do this)
  - Base pointer **RBP** is used to store that stack position

**Callee-saved**





# Calling a function – Example

■ Internally, the function will also need:

- Registers: RBX, R10, R11, R12
- Two 8-byte values as local variables

Note:

```
subq $16, %rsp
```

creates uninitialized space on the stack


We could use

```
pushq $0
```

```
pushq $0
```

instead to create the same space initialized to zero  
(or any other value we need).

**Stays the same while the  
function executes!**




# Calling a function – Example

■ Internally, the function will also need:

- Registers: **RBX**, R10, R11, **R12**
- Two 8-byte values as local variables


**Callee-saved**



# Calling a function – Example

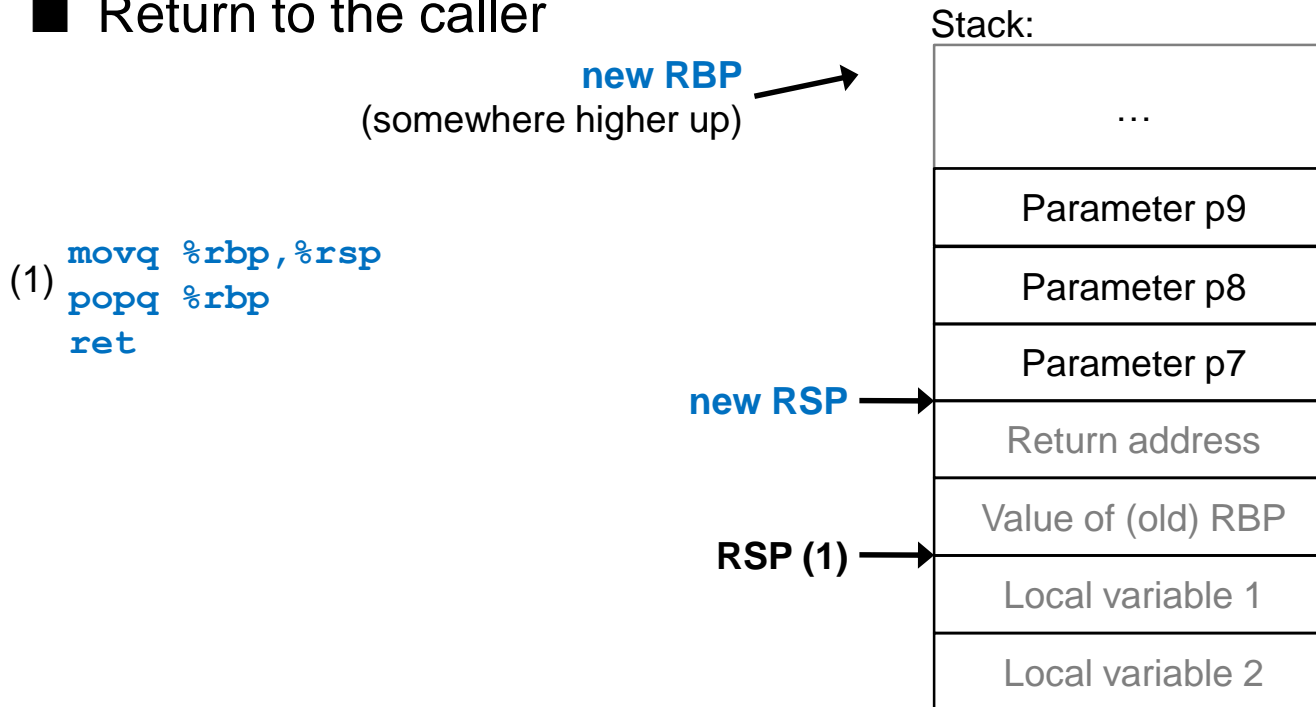
- At the end of the function, RAX is somehow set to the desired return value and the function has to clean up the stack
  - Restore saved Registers
    - In reverse order!

```
popq %r12  
popq %rbx
```



# Calling a function – Example

- Remove all local variables
  - Doesn't matter how many there are:  $RSP := RBP$  always removes all
- Restore the old RBP
- Return to the caller




# Calling a function – Example (callee)

doSomething:

<code>pushq %rbp</code>	<code># Store old base pointer</code>	}	Prologue
<code>movq %rsp,%rbp</code>	<code># Create new base pointer</code>		
<code>subq \$16,%rsp</code>	<code># Reserve space for 2 local variables</code>		
<code>pushq %rbx</code>	<code># Save old value on stack</code>		
<code>pushq %r12</code>	<code># R10 and R11 are caller-save!</code>		
<code>...</code>			
<code>movq 16(%rbp),%r12</code>	<code># Access parameter 7</code>		
<code>movq %r12,-16(%rbp)</code>	<code># Store it in local variable 2</code>		
<code>...</code>			
<code>movq %rdi,%rax</code>	<code># Set return value</code>		
<code>...</code>			
<code>addq \$10,%r10</code>	<code># Change the registers we "use"</code>		
<code>addq \$10,%r11</code>			
<code>addq \$10,%r12</code>			
<code>addq \$10,%rbx</code>			
<code>...</code>			
<code>popq %r12</code>	<code># Restore old register values</code>	}	Epilogue
<code>popq %rbx</code>			
<code>movq %rbp,%rsp</code>	<code># Destroy local variables</code>		
<code>popq %rbp</code>	<code># Restore old base pointer</code>		
<code>ret</code>	<code># Return to calling function</code>		

# Complete stack of example program



# Stack of example program in ddd

The screenshot shows the Data Display Debugger (DDD) interface. The main window displays the assembly code for a function named `doSomething`. The code includes instructions for pushing and popping registers, creating a new base pointer, reserving space for local variables, and returning to the calling function. A memory dump is visible on the left, showing the stack contents. Two dialog boxes are open: 'DDD: Examine Memory' and 'DDD: Registers'.

**DDD: Examine Memory**

Examine: 14, hex, giants (8), from: \$rsp

Print, Display, Close, Help

**DDD: Registers**

Registers	Value	Comment
rax	0x1	1
rbx	0x17	23
rcx	0x4	4
rdx	0x3	3
rsi	0x2	2
rdi	0x1	1
rbp	0x7fffffff0d0	0x7fffffff0d0
rsp	0x7fffffff0b0	0x7fffffff0b0
r8	0x5	5
r9	0x6	6
r10	0x14	20
r11	0x15	21
r12	0x11	17

Integer registers, All registers

Close, Help

**doSomething:**

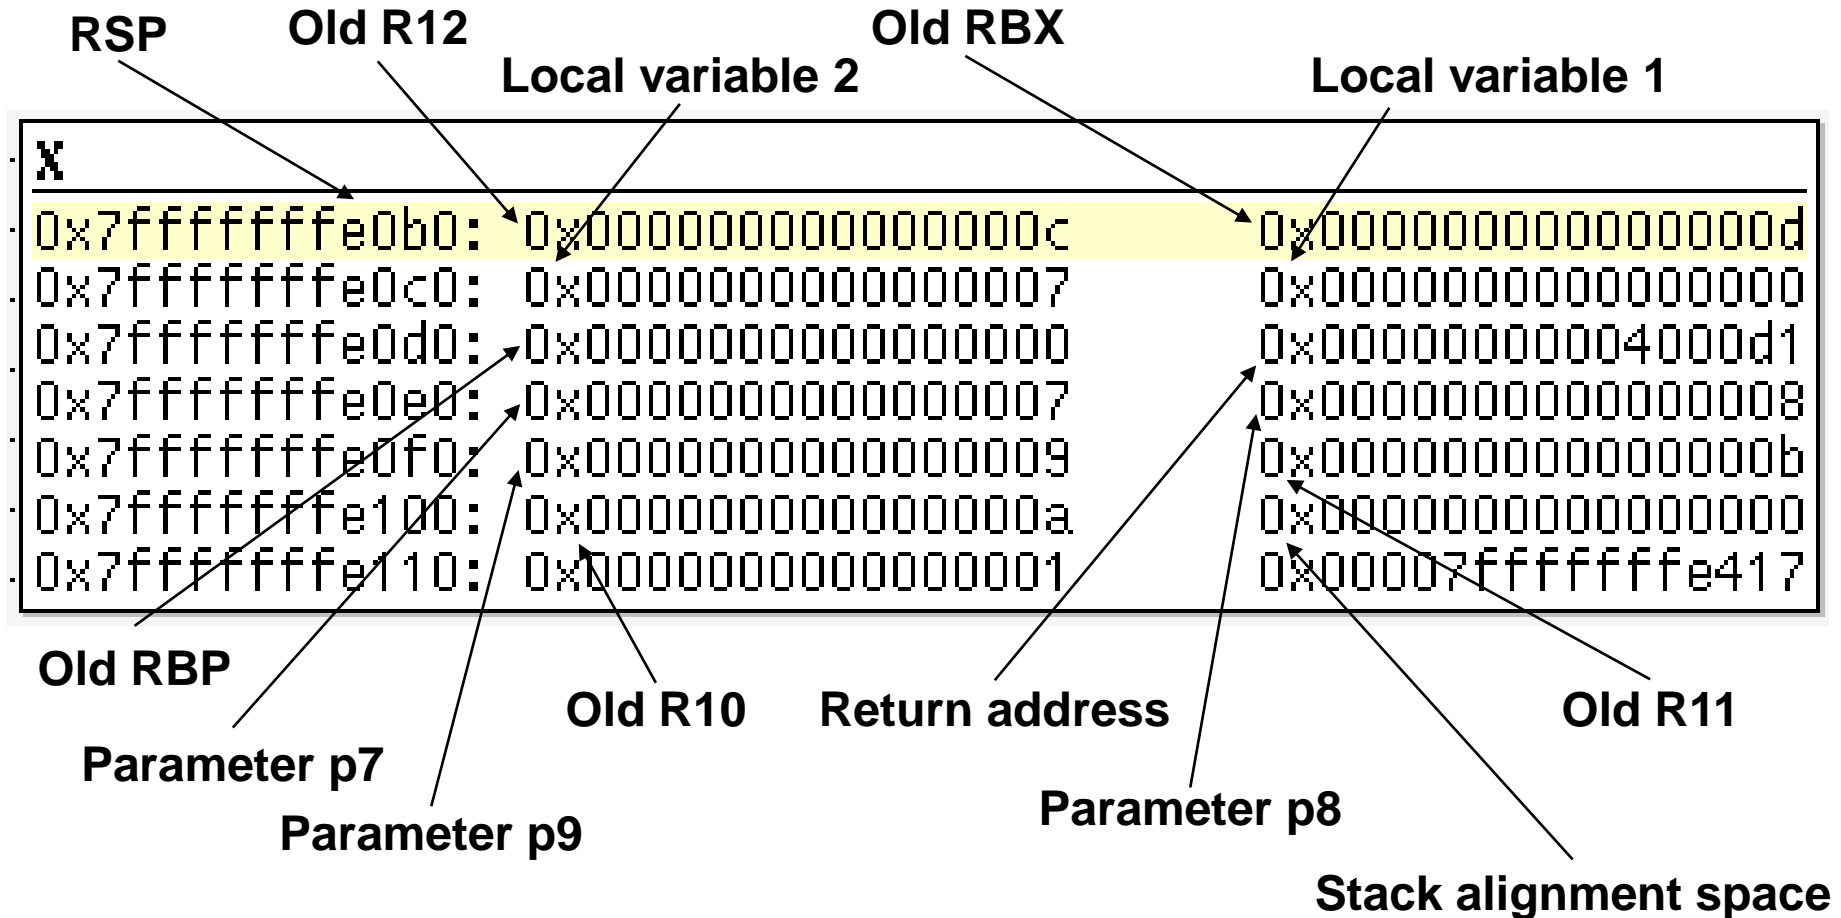
```
pushq %rbp          # Store old base pointer
movq %rsp,%rbp      # Create new base pointer
subq $16,%rsp       # Reserve space for 2 local variable
pushq %rbx          # Save old value on stack
pushq %r12          # R10 and R11 are caller-save!
# ...
movq 16(%rbp),%r12   # Access parameter 7
movq %r12,-16(%rbp) # Store it in local variable 2
# ...
movq %rdi,%rax       # Set return value
# ...
addq $10,%r10        # Change the registers we "use"
addq $10,%r11
addq $10,%r12
addq $10,%rbx
# ...
popq %r12            # Restore old register values
popq %rbx
movq %rbp,%rsp       # Destroy local variables
popq %rbp            # Restore old base pointer
ret                  # Return to calling function
```

(gdb) run  
Starting program: /mnt/function

Breakpoint 1, doSomething () at function.s:65  
(gdb) |

Display -1: `x /14xg \$rsp` (enabled)


# Stack of example program analyzed






# Calling a function – Example (callee)

```
doSomething:
    pushq %rbp          # Store old base pointer
    movq %rsp,%rbp      # Create new base pointer
                        # No need for RSP adjust., as less than 128 bytes
    pushq %rbx          # Save old value on stack
    pushq %r12          # R10 and R11 are caller-save!
    ...
    movq 16(%rbp),%r12   # Access parameter 7
    movq %r12,-16(%rsp)  # Store it in local variable 2
    ...
    movq %rdi,%rax      # Set return value
    ...
    popq %r12           # Restore old register values
    popq %rbx
    movq %rbp,%rsp      # Reset stack pointer always, even if unnecessary!
    popq %rbp           # Restore old base pointer
    ret                 # Return to calling function
```



- Variation: The function does not use explicit local variables, but uses the red zone (max. 128 bytes below RSP) instead
  - Still resets the base pointer

# Stack of example program (variant)



# Stack of red zone variant in ddd

Applications ▾ Places ▾ Data Display Debugger ▾ de ▾ Wed 16:02

DDD: /mnt/function-red.s

File Edit View Program Commands Status Source Data Help

Q: function-red.s:65

Lookup Findv Clear Watch Print Display Plot Hide Rotate Set Undisp

DDD: Examine Memory

Examine 16 hex giants (8) from \$rsp-32

Print Display Close Help

doSomething:

```

pushq %rbp          # Store old base pointer
movq %rsp,%rbp      # Create new base pointer
# No need for RSP adjustment, as less than 128 bytes
pushq %rbx          # Save old value on stack
pushq %r12           # R10 and R11 are caller-save!
# ...
movq 16(%rbp),%r12   # Access parameter 7
movq %r12,-16(%rsp)  # Store it in local variable 2
# ...
movq %rdi,%rax       # Set return value
# ...
addq $10,%r10        # Change the registers we "use"
addq $10,%r11
addq $10,%r12
addq $10,%rbx
# ...
popq %r12            # Restore old register values
popq %rbx
movq %rbp,%rsp       # Reset stack pointer always, even if not needed
popq %rbp            # Restore old base pointer
ret                  # Return to calling function

```

DDD: Registers

Registers	Value	Comment
rax	0x1	1
rbx	0x17	23
rcx	0x4	4
rdx	0x3	3
rsi	0x2	2
rdi	0x1	1
rbp	0x7fffffff0c0	0x7fffffff0c0
rsp	0x7fffffff0b0	0x7fffffff0b0
r8	0x5	5
r9	0x6	6
r10	0x14	20
r11	0x15	21
r12	0x11	17

Integer registers All registers

Close Help

DDD

Run Interrupt

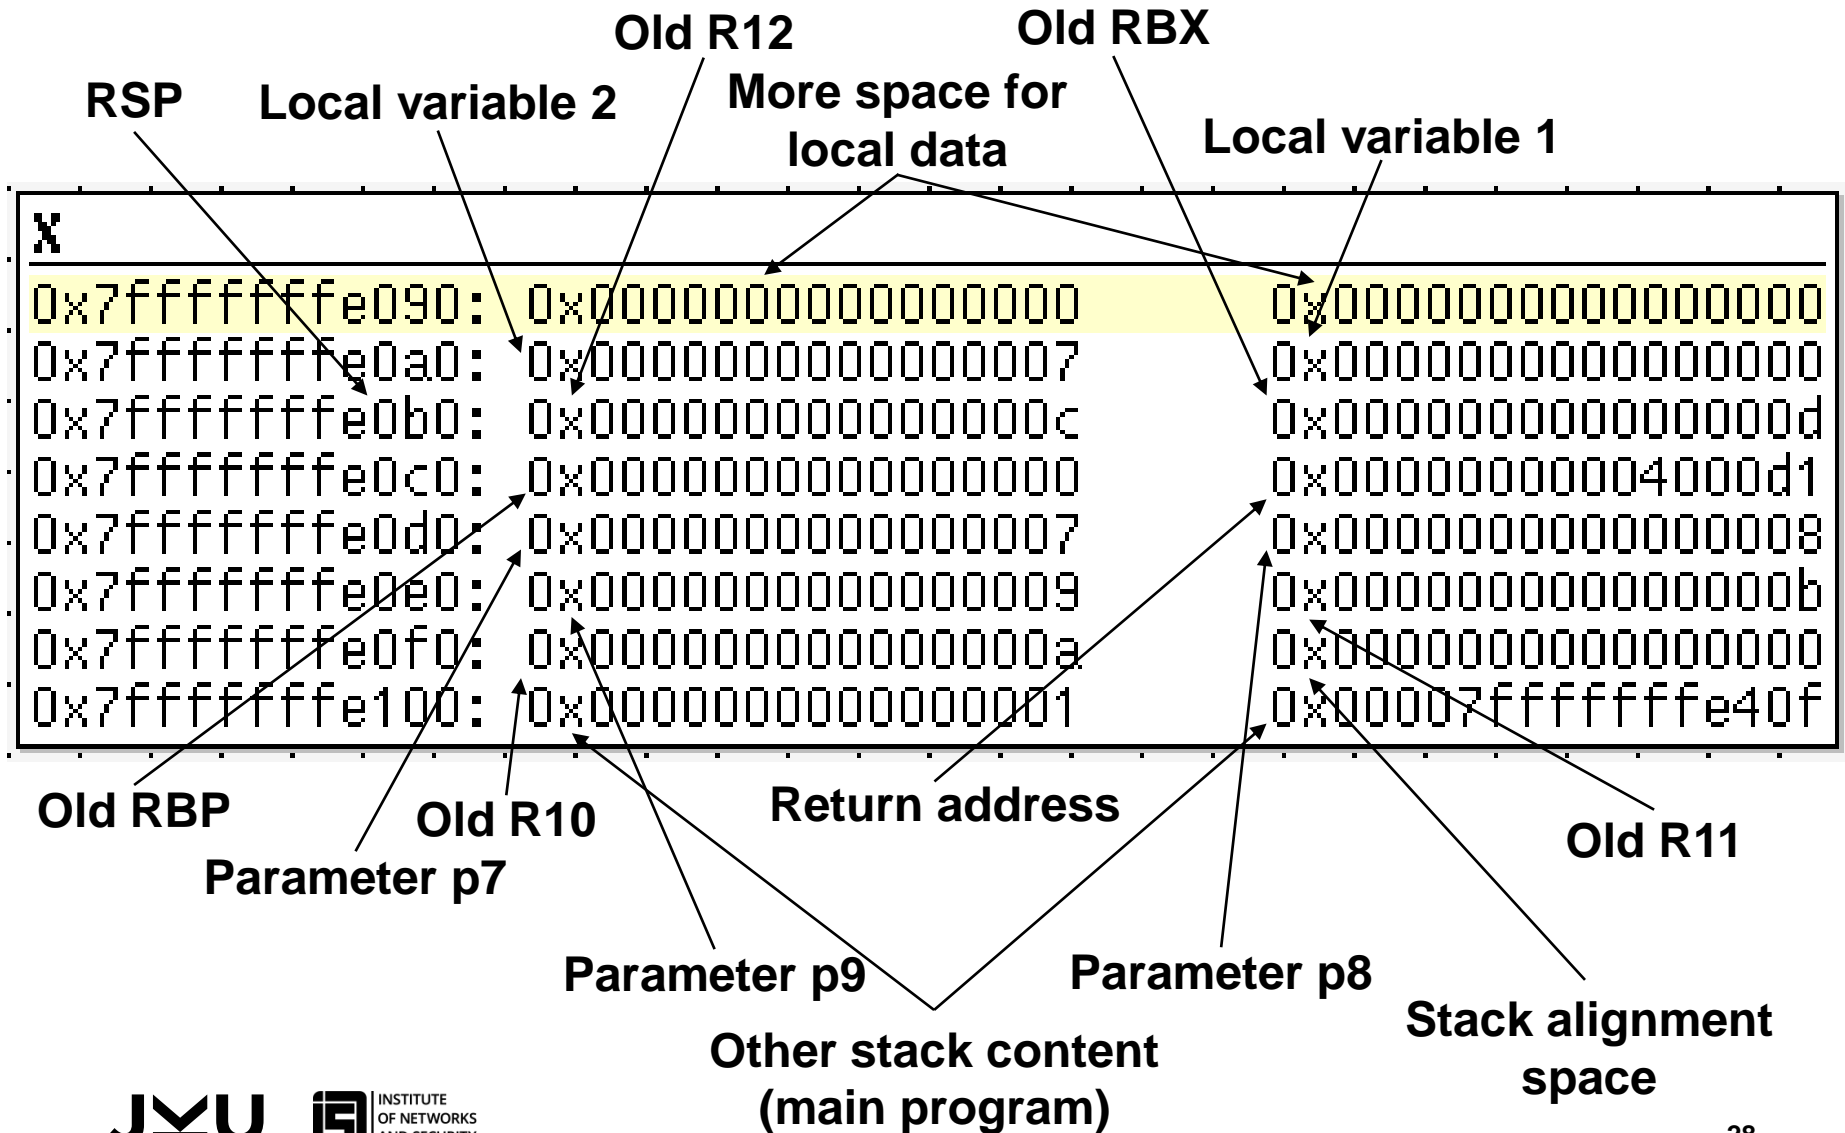
Step StepI Next NextI Until Finish Cont Kill Up Down Undo Redo Edit Make

(gdb) run  
Starting program: /mnt/function-red

Breakpoint 1, doSomething () at function-red.s:65  
(gdb)

Display -1: `x /16xg \$rsp-32` (enabled)

# Stack of red zone variant analyzed



# power1.s


```
# PURPOSE: Program to illustrate how functions work
#          This program will compute the value of 2^3 + 5^2
...
_start:
...
    movq $2,%rdi          # Store first argument
    movq $3,%rsi          # Store second argument
    call power             # Call the function
    movq %rax,%r12         # Save first result into temporary register

    movq $5,%rdi          # Store first argument
    movq $2,%rsi          # Store second argument
    call power             # Call the function
    movq %rax,%rdi        # Save second result into temporary register
    addq %r12,%rdi         # The second result is in %r12
                          # Add the first one and store in %rdi

    movq $60,%rax         # Exit (%rdi is returned)
    syscall
```

# power1.s

```
.type power, @function
power:
    pushq %rbp                # Save old base pointer
    movq  %rsp, %rbp          # Make stack pointer to %rbp
    subq  $8, %rsp            # Get room for our local storage
    pushq %rbx                # Preserve callee-save register
    movq  %rdi, %rbx          # Put first argument in %rbx
    movq  %rsi, %rcx          # Put second argument in %rcx
    movq  %rbx, -8(%rbp)       # Store current result
power_loop_start:
    cmpq  $1, %rcx            # If the power is 1, we are done
    je    end_power
    movq  -8(%rbp), %rax       # Move the current result into %rax
    imulq %rbx, %rax          # Multiply the current result by the base number
    movq  %rax, -8(%rbp)       # Store the current result
    decq  %rcx                # Decrease the power
    jmp   power_loop_start    # Run for the next power
end_power:
    movq  -8(%rbp), %rax       # Return value goes in %rax
    popq  %rbx                # Restore callee-save registers
    movq  %rbp, %rsp          # Restore the stack pointer
    popq  %rbp                # Restore the base pointer
    ret                       # Return to caller
```



# Notes on power1.s

## ■ `.type power,@function`

- Tells the linker that `power` should be treated as a function

## ■ Difference between `jmp` and `call`

- `jmp` modifies the RIP register to point to the new code location
- `call` additionally pushes the return address on the stack

## ■ The algorithm uses a local variable to **temporarily store** the result

- Also a register would be possible (if available, e.g. R12)
- But a register is not possible if the function calls another function and wants to pass a pointer to this variable as a parameter, as there is **no pointer to a register**
  - Registers do not have memory addresses!

## ■ This program does not work if the parameter `power` is zero

- See improved version `power2.s` in later slides

# power2.s

```
power:
    movq    $1,%rax
    cmpq    $0,%rsi          # If the power is 0, we return 1
    je      end_power
    movq    %rdi,%rax        # Prepare local variable for first round
power_loop_start:
    cmpq    $1,%rsi          # If the power is 1, we are done
    je      end_power
    imulq   %rdi,%rax         # Multiply the current result by the base number
    decq    %rsi              # Decrease the power
    jmp     power_loop_start  # Run for the next power
end_power:
    ret                    # Return to caller
```

- Optimized version: As this is a “leaf function” (it does not call any other functions itself), we can skip everything about the stack
  - No prologue, no epilogue → Sole stack content is return address
- We do not use any callee-safe registers, so we don’t have to save anything on the stack either
- Additionally check for “exponent 0” and return 1



# Factorial – Recursion example

## ■ Factorial of a number n

- Product of all numbers between 1 and number n
- Factorial of 7 =  $1 * 2 * 3 * 4 * 5 * 6 * 7$


## ■ Observation

- Factorial of 7 = factorial of 6 \* 7
- Generalized: **fac(n) = fac(n – 1) \* n**
- Base case: **fac(1) = 1**

## ■ Recursive definition

## ■ Implementation as a **recursive function**

- Function calls itself
- Returns when it reaches the base case



# factorial.s

```
.section .text
.globl _start
.globl factorial # this is not needed unless we want to share
                  # this function among other programs

_start:
    movq    $4,%rdi    # The factorial takes one argument - the
                       # number we want a factorial of (4 -> 24).

    call    factorial  # run the factorial function
    movq    %rax,%rdi  # factorial returns the answer in %rax, but
                       # we want it in %rdi to send it as our exit status

    movq    $60,%rax   # call the kernel's exit function
    syscall

.type factorial,@function
factorial:
    pushq   %rbp       # standard function stuff - we have to
                       # restore %rbp to its prior state before
                       # returning, so we have to push it

    movq    %rsp,%rbp  # This is because we don't want to modify
                       # the stack pointer, so we use %rbp.

    pushq   %rbx       # Save RBX (used for multiplication)
                       # Note: We could easily use e.g. R11 to
                       # avoid needing the stack!
```

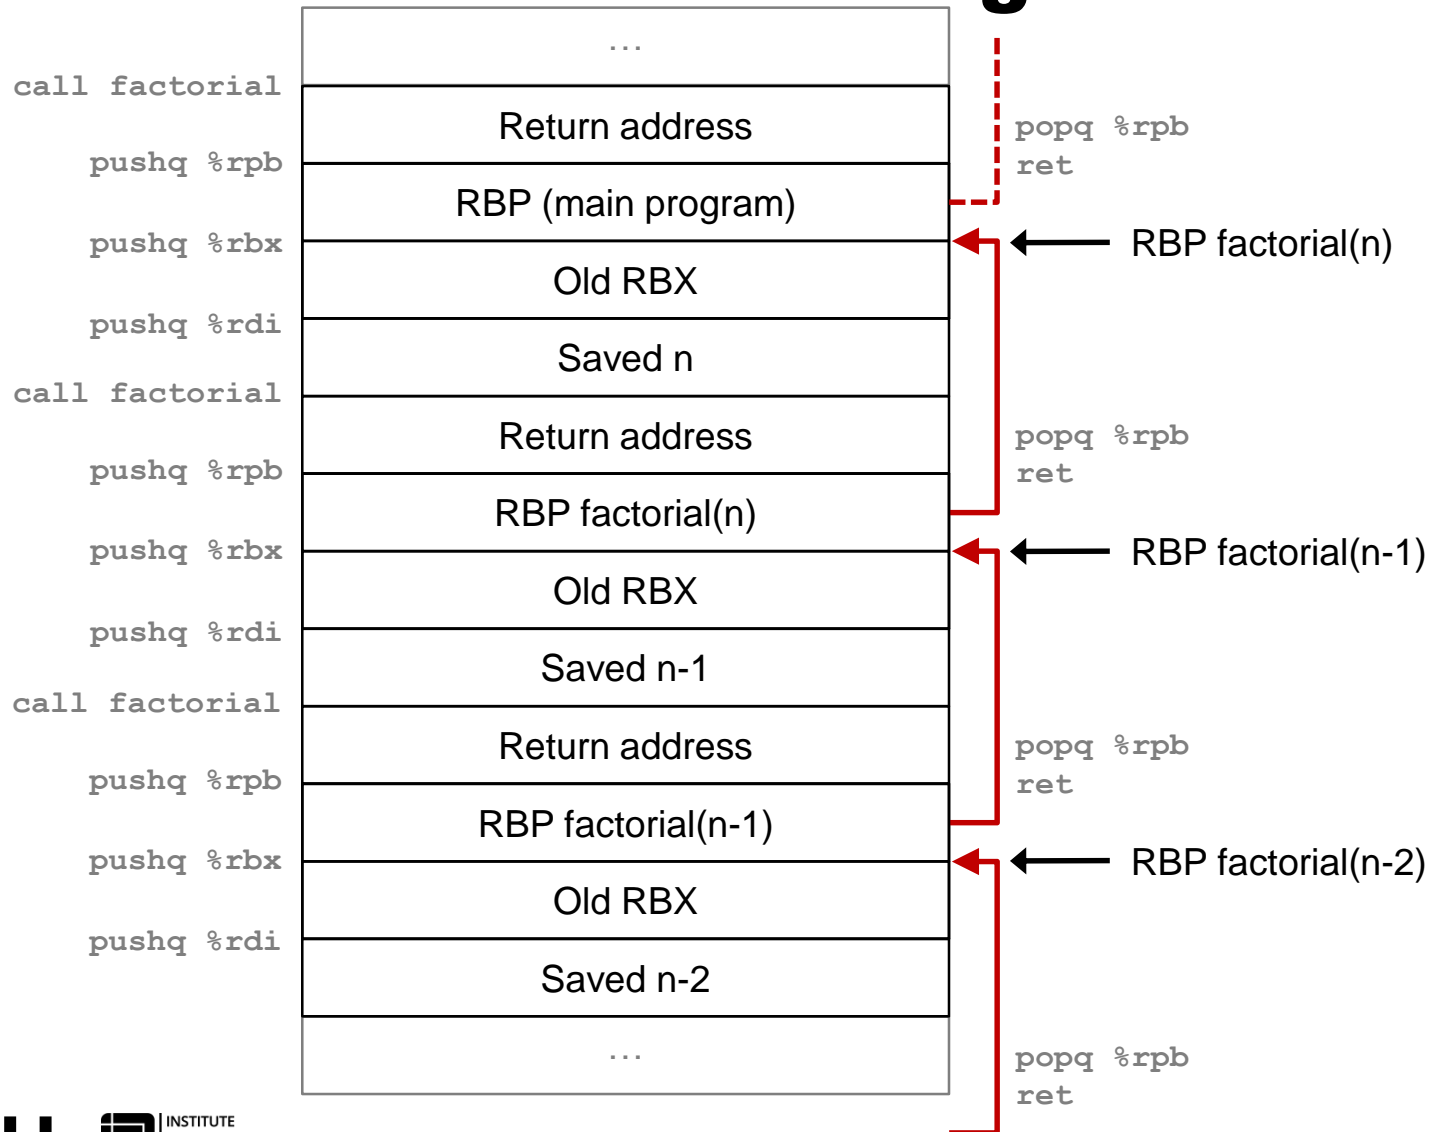
# factorial.s

```
check_base_case0:
    movq    $1,%rax
    cmpq    $0,%rdi    # If the number is 0, we return 1
    je     end_factorial
check_base_case1:
    cmpq    $1,%rdi    # If the number is 1, that is our base
    je     end_factorial # case, and we simply return (1 is
                        # already in %rax as the return value)

    pushq   %rdi        # save our own parameter for later
    decq    %rdi        # decrease the value
    call    factorial   # call factorial
    popq    %rbx        # retrieve our own parameter
    imulq   %rbx,%rax   # multiply it by the result of the last
                        # call to factorial (in %rax); the answer
                        # is stored in %rax, which is good since
                        # that's where return values go.

end_factorial:
    popq    %rbx        # restore old value
    movq    %rbp,%rsp   # standard function return stuff - we
    popq    %rbp        # have to restore %rbp and %rsp to where
                        # they were before the function started
    ret      # return from the function
```

# Stack of factorial.s during recursion



# binom.s


- Binomial coefficient “n over k”
  - $(n \text{ over } 0) = (n \text{ over } n) = 1$  (base case)
  - $(n \text{ over } k) = (n - 1 \text{ over } k - 1) + (n - 1 \text{ over } k)$  (recursive case)
- Function binom(n, k)
  - $\text{binom}(n, 0) = \text{binom}(n, n) = 1$
  - $\text{binom}(n, k) = \text{binom}(n-1, k-1) + \text{binom}(n-1, k)$
- Differences to factorial.s
  - 2 base cases
  - 2 recursive calls in general case
  - Need to save intermediate result of first call
- Note: Code does not check the parameters for validity
- Note: Return address + RBP → Stack is again correctly aligned

# binom.s

```
.type binom,@function
                                # RDI = n, RSI = k
binom:
    pushq %rbp                  # standard function
                                # restore %rbp to its original value
                                # returning, so we need to push it
    movq  %rsp,%rbp             # This is because we need to push the stack pointer
                                # 8(%rbp) holds the return address
    subq  $32,%rsp              # get room for local n, local k and
                                # result of first recursive call
                                # Additional 8 Bytes for stack alignment
                                # in recursive calls

check_base_case1:
    cmpq  $0,%rsi               # If k is 0, we return 1
    jne  check_base_case2
    movq  $1,%rax
    jmp  end_binom

check_base_case2:
    cmpq  %rdi,%rsi             # If n = k, we return 1
    jne  general_case
    movq  $1,%rax
    jmp  end_binom
```




# binom.s

general\_case:

```
# Note: Parameters are passed in registers
# so we do not have a "backup copy" on the stack
movq %rdi, -8(%rbp) # save n
movq %rsi, -16(%rbp) # save k
decq %rdi           # decrease n
decq %rsi           # decrease k
# first recursive call: (n - 1 over k - 1)
call binom          # recursive call
movq %rax, -24(%rbp) # save value of first recursive call
movq -8(%rbp), %rdi  # restore n
movq -16(%rbp), %rsi # restore k
decq %rdi           # decrease n
# second recursive call: (n - 1 over k)
call binom          # recursive call
# %rax holds result of second recursive call
addq -24(%rbp), %rax # compute sum of recursive calls = result
# %rax holds result
```

end\_binom:

```
movq %rbp, %rsp      # standard function return stuff - we
popq %rbp            # have to restore %ebp and %esp to where
                    # they were before the function started
ret                  # return from the function
```



# THANK YOU FOR YOUR ATTENTION!

**Slides by: Michael Roland and Michael Sonntag**

michael.roland@ins.jku.at, michael.sonntag@ins.jku.at

+43 (732) 2468 - 4137

S3 235 (Science park 3, 2<sup>nd</sup> floor)



JOHANNES KEPLER  
UNIVERSITÄT LINZ



INSTITUTE  
OF NETWORKS  
AND SECURITY

<http://www.ins.jku.at>



**JOHANNES KEPLER  
UNIVERSITÄT LINZ**

Altenberger Straße 69  
4040 Linz, Österreich  
[www.jku.at](http://www.jku.at)