

Files



Systems Programming



Michael Sonntag
Institute of Networks and Security



The File concept of Unix

- Files are accessed as a sequential stream of bytes
- Opening a file returns a file descriptor (which is a number)
 - Might also be a pointer – but for the user it is **completely opaque!**
 - **MUST** be retained – access to a file is **ONLY** possible through this!
- File permissions
 - Three modes: read, write, execute (rwx)
 - (110) means permission to read and write but not to execute
 - Three different sets
 - user: every file has an owner
 - group: every file belongs to a single group (of arbitrary users)
 - others: users that are neither owner nor in the group
 - Octal encoding
 - Convert binary permission sets to octal numbers
 - Combine them to one number and prefix with 0
 - 0754 = owner can read, write & execute the file,
group can read and execute,
everyone else can only read

user	group	other
rwx	r-x	r--
111	101	100
7	5	4

Dealing with Files

- Lifecycle of all files:
 - ☐ Open file & check for success
 - ☐ Read from / write to file & check for success every time
 - ☐ Close file
 - & check for success → But this is rare and there is little you can do
- All files are closed when the program terminates
 - ☐ Unfortunately, there is little guarantee how → data loss possible!
- Opening a file is not possible yourself → Ask the OS to do it for you
 - ☐ For this you need the file name (including the path - or it will use the current directory of the program)
 - This course: statically defined (data section) or program parameter
- Reading and writing is not possible directly, i.e. from a register
 - ☐ You can only read to / write from memory
 - ☐ So we need a buffer there

Opening a file

- RAX: 2 (=sys_open system call)
- RDI: Address of filename (must be nul-terminated, i.e. a C string)
- RSI: Flags (read, write, read&write, append...)
 - ☐ Must contain one of O_RDONLY, O_WRONLY, or O_RDWR
 - ☐ 0..N creation flags O_CREAT, O_TMPFILE, O_TRUNC...
 - ☐ 0..N status flags: O_APPEND, O_ASYNC...
- RDX: Mode = Permissions for file (when creating one)
 - ☐ Use 0777 at most (we use 0666)
 - Linux: Also 04000=SUID, 02000=SGID, 01000=Sticky are possible
- Return value RAX: File descriptor
 - ☐ ≥ 0 : Success – file descriptor
 - ☐ < 0 : Error – Negative of error number
 - Example: -13 = Error number 13 = EACCESS = Access not allowed
- See also <http://man7.org/linux/man-pages/man2/open.2.html>

Reading from a file

- RAX: 0 (=sys_read)
- RDI: File descriptor
- RSI: Address of buffer to be filled with file data
 - Must contain space for RDX bytes!
 - This is a binary buffer, so there is no termination (nul or other)
- RDX: Number of bytes to read at most
 - The OS will **always try** to give you that much data, but there is no guarantee: file is not long enough, network problem...
- Return value RAX: Number of characters **actually read**
 - >0: Success – RAX bytes placed in buffer
 - =0: Success – End of file reached (& no data read)
 - No data available, but not EOF → Call blocks (or returns error number EAGAIN; see option O_NONBLOCK)!
 - <0: Error – Negative of error number

Writing to a file

- RAX: 1 (=sys_write)
- RDI: File descriptor
- RSI: Address of buffer with data to be written to file
 - Must contain at least RDX bytes!
- RDX: Number of bytes to write
 - Note: The OS will **always try** to write the full amount, but there is no guarantee: disk full, network problem...
- Return value RAX: Number of characters **actually written**
 - ≥ 0 : Success – RAX bytes written to file
 - But not necessarily yet on disk – might be in OS buffer only!
 - Might also block if O_NONBLOCK is not set
 - < 0 : Error – Negative of error number

Closing a file

- RAX: 3 (=sys_close)
- RDI: File descriptor
- Return value RAX
 - =0: Success
 - <0: Error – Negative of error number
- Note: Writing to a file on a network filesystem might report writing errors only on closing the file (but not on the individual write, as storing the data in the local buffer succeeds!)
- Note: Closing a file is no guarantee that the data is on the disk
 - Use fsync before (RAX 75, RDI file handle), but this may block
 - Note: Guarantees that all file data was sent to the device. This still is no guarantee it is permanently stored (internal buffers)!
 - Note: No guarantees about the file entry (=directory content)

System calls for file manipulation

System call	RAX (cmd.)	RDI (parameter 1)	RSI (parameter 2)	RDX (parameter 3)	RAX (return value)
SYS_OPEN	2	Pointer to filename	Flags (O_RDONLY, ...)	Create mode (e.g. 0666)	File descriptor or error number
SYS_READ	0	File descriptor	Pointer to data buffer	Max. number of bytes to read	Actual number of bytes read or error number
SYS_WRITE	1	File descriptor	Pointer to data buffer	Number of bytes to write	Actual number of bytes written or error number
SYS_CLOSE	3	File descriptor	---	---	0 (success) or error number

Buffers – Space for data

- Buffers must be reserved “somehow”:
 - Static: Define in assembler file
 - Stack: Reduce RSP
 - Not recommended except for very small buffers
 - Heap: Explicit memory reservation (see later)
 - Recommended for large buffers
- Static buffers: Declare in section BSS
 - BSS: On many (=not all!) systems initialized to all zeros

```
.section .bss
.lcomm my_buffer, 500      # Create a symbol for the start address
                           # Note: No „$“ for the length!

.....

# RDI already contains the file descriptor
movq $0, %rax              # Read from file into buffer
movq $my_buffer, %rsi      # Store start address of buffer
movq $500, %rdx            # Store length of buffer
syscall                   # Read from file; result in RAX
```

Standard file descriptors

- Three file descriptors are already open per default

- STDIN

- Represents input read from keyboard
 - End of input → press <CTRL-d>
 - File descriptor 0

- STDOUT

- Represents output written to screen
 - File descriptor 1

- STDERR

- Represents error output written to screen
 - File descriptor 2

- Do NOT close them; they cannot be reopened!

- Unless you really know what you (want to) do...

- E.g. for daemons/services running only in the background

Unix file paradigm

- The default behavior of most UNIX programs is to
 - ☐ Read input from standard input (STDIN)
 - ☐ Write output to standard output (STDOUT)
 - ☐ Write error output to standard error (STDERR)

- The paradigm of UNIX is to treat all input/output systems as files
 - ☐ Network connections
 - ☐ Serial port
 - ☐ Audio devices
 - ☐ Harddisks
 - ☐ etc.


Redirecting Input/Output

- Redirect stdin to file (=read input from file as if it was typed in manually)
 - `sort < list.input`
- Redirect stdout to file (= output is written into file **instead** of printed on screen)
 - `ls > ls.output`
- Redirect stderr to file
 - `ls 2> ls.error`
- Redirect stderr to stdout
 - `ls 2>&1`
- Redirect stdout to stderr
 - `ls 1>&2`
- Redirect stdout and stderr to file
 - `ls &> ls.output`
- Redirect and append stdout to file
 - `ls >> ls.output`
- Redirect and append stderr to file
 - `ls 2>> ls.error`

Pipes

■ Pipes

- They connect programs, similar to a physical pipeline
- Feed output from program A directly as input to program B
 - Connects STDOUT from first program to STDIN from second program
- Often | used as pipe symbol
- `cat file.txt | sort | uniq`



Comparing files

- Comparing files (e.g. assignment exemplary output):
 - Command-line only (or you would need to find other tools)
 - Such exist for all OS, but they are mostly for much more complicated tasks!
 - `cmp -b your_file exemplary_file`
 - Return value: 0 = identical, 1 = different
 - -b also prints the differing bytes
 - `diff -u your_file exemplary_file`
 - -u also shows the “surrounding” – a few lines before and after

Comparing files - Example

■ `diff -u maximum.s maximum_new.s`

--- maximum.s 2017-10-18 15:39:14.000000000 +0200

+++ maximum_new.s 2019-03-26 13:51:15.546264824 +0100

@@ -5,7 +5,7 @@

#VARIABLES: The registers have the following uses:

#

%rdx - Holds the index of the data item being examined

- *# %rdi - Largest data item found*

+ *# %rdi - Largest data item found until now*

Changed line

%rax - Current data item

#

The following memory locations are used:

@@ -24,11 +24,11 @@

.globl _start

_start:

movq \$0, %rdx # move 0 into the index register

- *movq data_items(,%rdx,8), %rax # load the first byte of data*

Removed line

movq %rax, %rdi # since this is the first item, %rax is

the biggest

start_loop: # start loop

+ *movq data_items(,%rdx,8), %rax # load the first byte of data*

Inserted line

cmpq \$0, %rax # check to see if we've hit the end

je loop_exit

incq %rdx # load next value

THANK YOU FOR YOUR ATTENTION!

Slides by: Michael Sonntag
michael.sonntag@ins.jku.at
+43 (732) 2468 - 4137
S3 235 (Science park 3, 2nd floor)